

# PROBABILISTIC & LANGUAGE MODELS

---

## Chapter 7

# Outline

- Probability ranking principle
- Classical probabilistic model
  - Binary Independence Model
  - BM25
  - feedback methods
- Language Models

# Probability Ranking Principle

- Robertson (1977)
  - “If a reference retrieval system’s response to each request is a *ranking* of the documents in the collection in order of decreasing *probability of relevance* to the user who submitted the request,
  - where the *probabilities* are *estimated* as *accurately* as possible on the basis of whatever data have been made available to the system for this purpose,
  - the overall *effectiveness* of the system to its user will be the *best* that is obtainable on the basis of those data.”
- Basis for most probabilistic approaches to IR

# Let's dissect the PR

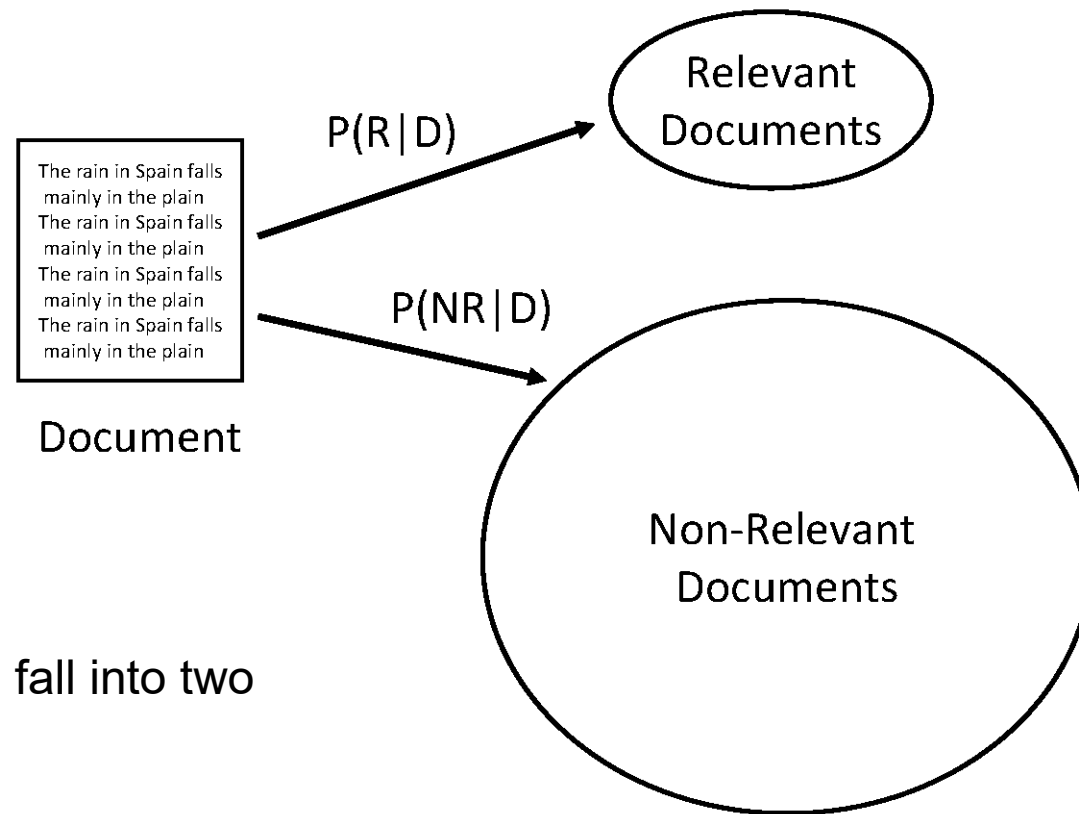
- rank documents ... by probability of relevance
  - $P(\text{relevant} \mid \text{document})$  or  $P(R \mid D)$
- estimated as accurately as possible
  - $P_{\text{est}}(R \mid D) \rightarrow P_{\text{true}}(R \mid D)$  in some way
- based on whatever data is available to system
  - $P_{\text{est}}(R \mid D, \text{query}, \text{context}, \text{user profile}, \dots)$
- best possible accuracy one can achieve with that data
  - recipe for a perfect IR system: just need  $P_{\text{est}}(R \mid \dots)$
  - strong stuff, can this really be true?

# Probability of relevance

- What is:  $P_{\text{true}}(\text{relevant} \mid \text{doc}, \text{qry}, \text{user}, \text{context})$  ?
  - isn't relevance just the user's opinion?
  - user decides relevant or not, what's the “probability” thing?
- “user” does not mean the human being
  - doc, qry, user, context ... representations
    - parts of the real thing that are available to the system
  - typical case:  $P_{\text{true}}(R \mid D, \text{query})$ 
    - query: 2-3 keywords, user profile unknown, context not available
    - whether document is relevant is uncertain
      - depends on the factors which are not available to our system
    - think of  $P_{\text{true}}(R \mid D, \text{qry})$  as proportion of all unseen users/contexts/... for which the document would have been judged relevant

Take away ... this is a difficult problem!

# IR as Classification



For a given query, documents fall into two classes

- relevant (R) and non-relevant (NR)
- compute  $P(R|D)$  and  $P(NR|D)$ 
  - retrieve if  $P(R|D) > P(NR|D)$

How do we compute these probabilities?


# Generative vs. Discriminative Classifiers

Extra Slides

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Discriminative classifiers (also called ‘informative’ by Rubinstein&Hastie):

1. Assume some functional form for  $P(Y|X)$
2. Estimate parameters of  $P(Y|X)$  directly from training data




# Generative vs. Discriminative Classifiers

Extra Slides

Training classifiers involves estimating  $f: X \rightarrow Y$ , or  $P(Y|X)$

Generative classifiers

1. Assume some functional form for  $P(X|Y)$ ,  $P(X)$
2. Estimate parameters of  $P(X|Y)$ ,  $P(X)$  directly from training data
3. Model predictions based on  $P(Y|X = x_i)$





# Bayes Classifier

- Bayes Decision Rule
  - A document  $D$  is relevant if  $P(R|D) > P(NR|D)$
- Estimating probabilities
  - use Bayes Rule
  - classify a document as relevant if  $P(R|D) > P(NR|D)$ .

$$P(R|D) = \frac{P(D|R)P(R)}{P(D)}$$

- *likelihood ratio*

- *In practice, we rank documents by*

$$\frac{P(D|R)}{P(D|NR)} > \frac{P(NR)}{P(R)}$$

Now the question is how to compute  $P(D|R)$  and  $P(D|NR)$ .

# Maximum A Posteriori (MAP) Classifier

- Given the data feature vector  $\mathbf{x}$ , we would like to find the class with the largest probability:

$$\hat{C} \triangleq \arg \max_C p(C|\mathbf{x})$$

- To accomplish this, we iterate through all possible classes  $C_1, C_2, \dots, C_K$  and evaluate the quantity  $p(C_i|\mathbf{x})$ , and pick the class that has the largest probability.

# MAP Classifier: Bayes Rule (1/4)

- How do we compute  $p(C_i|\mathbf{x})$ ? Apply Bayes' Rule!

$$p(A|B)p(B) = p(B|A)p(A) \Rightarrow p(A|B) = \frac{p(B|A)p(A)}{p(B)}$$

likelihood  $\swarrow$  prior  $\swarrow$   
 $\nwarrow$  evidence

- Applying Bayes' rule to  $p(C_i|\mathbf{x})$ , we obtain:

$$p(C_i|\mathbf{x}) = \frac{p(\mathbf{x}|C_i)p(C_i)}{p(\mathbf{x})}$$

## MAP Classifier: Bayes Rule (2/4)

$$p(C_i | \mathbf{x}) = \frac{p(\mathbf{x} | C_i) p(C_i)}{p(\mathbf{x})}$$

- $p(C_i)$  is called the prior probability of a class
- $p(\mathbf{x} | C_i)$  is called the likelihood of the data (what is the probability of observing  $\mathbf{x}$  if the class was  $C_i$ )
- $p(\mathbf{x})$  is the probability of seeing the data so its simply a normalizing factor (applies to all  $C_i$ ), so doesn't affect which  $C_i$  attains MAP.

# MAP Classifier: Bayes Rule (3/4)

- To compute the normalizing factor, we use:

$$p(\mathbf{x}) = \sum_{i=1}^K p(\mathbf{x}|C_i)p(C_i)$$

So that

$$\sum_{i=1}^K p(C_i|\mathbf{x}) = 1$$

Hence, 
$$p(C_i|\mathbf{x}) = \frac{p(\mathbf{x}|C_i)p(C_i)}{\sum_{\ell=1}^K p(\mathbf{x}|C_\ell)p(C_\ell)}$$

But really, we don't need to calculate the denominator  $p(\mathbf{x})$ , so we can simply write it as

$$(p(C_i|\mathbf{x}) = p(\mathbf{x}|C_i)p(C_i)$$

## MAP Classifier: Bayes Rule (4/4)

- Usually, the likelihood  $p(\mathbf{x}|C_i)$  is difficult to compute because it is N-dimensional (length of feature vector).
- This is because the distribution considers correlations between the features when computing the likelihood.
- We can write  $p(\mathbf{x}|C_i)$  equivalently as:

$$p(\mathbf{x}|C_i) = p(x_1, x_2, \dots, x_N|C_i)$$

which makes the dependence on individual features explicit.

# Naive Bayes Classifier

- So,  $P(c)$  is easy to compute...
- What about  $P(x_1, x_2, \dots, x_n \mid c)$  ?
  - The probability of the class given the features.
- This is hard to compute given that there are many features.
  - $O(|X|^n |C|)$  parameters
- Also, its hard to generate this unless we have a large training dataset.

# Naive Bayes Classifier

$$P(x_1, x_2, \dots, x_n | c)$$

- Bag of Words assumption : order of the words don't matter
- Conditional independence: Assume the feature probability  $P(x_i|c)$  are independent given the class  $c$ .

$$P(x_1, x_2, \dots, x_n | c) = P(x_1 | c) \cdot P(x_2 | c) \cdot \dots \cdot P(x_n | c)$$

Obviously this assumption may not always be true, but this simplification allows us to solve the more easily.



# Naïve Bayes Classifier: Independence Assumption

- The Naïve Bayes classifier introduces one major assumption regarding the features: independence
- That is, the Naïve Bayes classifier assumes:

$$p(\mathbf{x}|C_i) = p(x_1, x_2, \dots, x_N|C_i) = \prod_{n=1}^N p(x_n|C_i)$$

- That is, the complicated likelihood  $p(\mathbf{x}|C_i)$  can now be factored into a product of N 1-dimensional likelihoods, which are easy to compute.
- It is good to note that independence implies that the features are not correlated (but generally not the other way around---so independence is a stronger assumption).

# Binary independence model

- Objective: Need to compute  $P(D|R)$  and  $P(D|NR)$
- Assumptions
  - In this model, documents are represented as a *vector of binary features*
    - $D = (d_1, d_2, \dots, d_t)$ , where  $d_i = 1$  if term  $i$  is present in the document, and 0 otherwise.
  - Also assume term independence (also known as the Naïve Bayes assumption).
- Hence,  $P(D|R)$  is computed as the product of the individual term probabilities and similarly for  $P(D|NR)$ .

$$P(D|R) = \prod_{i=1}^t P(d_i|R)$$

# Binary Independence Model

$$\begin{aligned}
 \frac{P(D|R)}{P(D|NR)} &= \prod_{i:d_i=1} \frac{p_i}{s_i} \cdot \prod_{i:d_i=0} \frac{1-p_i}{1-s_i} \\
 &= \prod_{i:d_i=1} \frac{p_i}{s_i} \cdot \left( \prod_{i:d_i=1} \frac{1-s_i}{1-p_i} \cdot \prod_{i:d_i=1} \frac{1-p_i}{1-s_i} \right) \cdot \prod_{i:d_i=0} \frac{1-p_i}{1-s_i} \\
 &= \prod_{i:d_i=1} \frac{p_i(1-s_i)}{s_i(1-p_i)} \cdot \boxed{\prod_i \frac{1-p_i}{1-s_i}}
 \end{aligned}$$

Document-  
independent

- $p_i$  is probability that term  $i$  occurs in relevant document,
- $s_i$  is probability that term  $i$  occurs in non-relevant document

## Binary Independence Model

$$= \prod_{i:d_i=1} \frac{p_i(1-s_i)}{s_i(1-p_i)} \cdot \prod_i \frac{1-p_i}{1-s_i}$$

Second term is the same for all documents, so we can drop it for purpose of ranking.

- Scoring function is

$$\sum_{i:d_i=1} \log \frac{p_i(1-s_i)}{s_i(1-p_i)}$$

# Binary Independence Model

- Scoring function is  $\sum_{i:d_i=1} \log \frac{p_i(1-s_i)}{s_i(1-p_i)}$
- Query provides information about relevant documents
- If we assume  $p_i$  constant for terms in query,  $s_i$  approximated by entire collection, get *idf*-like weight

$$\log \frac{0.5(1 - \frac{n_i}{N})}{\frac{n_i}{N}(1 - 0.5)} = \log \frac{N - n_i}{n_i}$$

N number of documents in collection

$N_i$  is the number of documents that contain term  $i$

Similar to idf weight, where there is no tf component since documents have binary features

# Contingency Table

What if we had some kind of relevance feedback from users?

	Relevant	Non-relevant	Total
$d_i = 1$	$r_i$	$n_i - r_i$	$n_i$
$d_i = 0$	$R - r_i$	$N - n_i - R + r_i$	$N - r_i$
Total	$R$	$N - R$	$N$

$n_i$ : # of docs that contain term  $i$

$N$ : total # of docs in collection

$r_i$ : # of relevant docs containing term  $i$

$R$ : # relevant docs

# Contingency Table

	Relevant	Non-relevant	Total
$d_i = 1$	$r_i$	$n_i - r_i$	$n_i$
$d_i = 0$	$R - r_i$	$N - n_i - R + r_i$	$N - r_i$
Total	$R$	$N - R$	$N$

$$p_i \equiv \frac{r_i}{R}$$

The number of relevant documents that contain a term divided by the number of relevant documents)

$$s_i \equiv \frac{(n_i - r_i)}{(N - R)}$$

The number of non-relevant documents that contain a term divided by the total number of non-relevant documents)

## Binary Independence Model (Cont.)

$$p_i = (r_i + 0.5)/(R + 1)$$

$$s_i = (n_i - r_i + 0.5)/(N - R + 1)$$

Gives scoring function:

$$\sum_{i:d_i=q_i=1} \log \frac{(r_i+0.5)/(R-r_i+0.5)}{(n_i-r_i+0.5)/(N-n_i-R+r_i+0.5)}$$



# BM25 Ranking Algorithm

- BM25 is a scoring function based on the binary independence model
  - Stands for Best Match & 25 is a numbering scheme used by its director Robertson.
  - It includes document and query term weights.
- It has performed very well in TREC experiments, although it does not exactly a formal model.

# BM25

- $N$  : total number of documents in the collection
  - $R$  : number of relevant documents
  - $n_i$  : number of documents with term  $i$
  - $r_i$  : number of relevant documents with term  $i$
  - $f_i$  : frequency of term  $i$  in the document
  - $qf_i$  : frequency of term  $i$  in the query
  - $k_1, k_2, K$  : parameters whose values are set empirically
- } same as before

$$\sum_{i \in Q} \log \frac{(r_i + 0.5) / (R - r_i + 0.5)}{(n_i - r_i + 0.5) / (N - n_i - R + r_i + 0.5)} \cdot \frac{(k_1 + 1) f_i}{K + f_i} \cdot \frac{(k_2 + 1) q f_i}{k_2 + q f_i}$$

## BM25 Example

- Query with two terms, “president” & “lincoln”, ( $qf = 1$ )
- No relevance information ( $r$  and  $R$  are zero)
- $N = 500,000$  documents
- “*president*” occurs in 40,000 documents ( $n_1 = 40,000$ )
- “*lincoln*” occurs in 300 documents ( $n_2 = 300$ )
- “president” occurs 15 times in doc ( $f_1 = 15$ )
- “*lincoln*” occurs 25 times ( $f_2 = 25$ )
- document length is 90% of the average length ( $dl/avdl = .9$ )
- $k_1 = 1.2$ ,  $b = 0.75$ , and  $k_2 = 100$
- $K = 1.2 \cdot (0.25 + 0.75 \cdot 0.9) = 1.11$

# BM25 Example

$$BM25(Q, D) =$$

Log<sub>e</sub>() is  
used

$$\begin{aligned} & \log \frac{(0 + 0.5)/(0 - 0 + 0.5)}{(40000 - 0 + 0.5)/(500000 - 40000 - 0 + 0 + 0.5)} \\ & \times \frac{(1.2 + 1)15}{1.11 + 15} \times \frac{(100 + 1)1}{100 + 1} \\ & + \log \frac{(0 + 0.5)/(0 - 0 + 0.5)}{(300 - 0 + 0.5)/(500000 - 300 - 0 + 0 + 0.5)} \\ & \times \frac{(1.2 + 1)25}{1.11 + 25} \times \frac{(100 + 1)1}{100 + 1} \end{aligned}$$

$$\begin{aligned} &= \log 460000.5/40000.5 \cdot 33/16.11 \cdot 101/101 \\ &+ \log 499700.5/300.5 \cdot 55/26.11 \cdot 101/101 \\ &= 2.44 \cdot 2.05 \cdot 1 + 7.42 \cdot 2.11 \cdot 1 \\ &= 5.00 + 15.66 = 20.66 \end{aligned}$$

## BM25 Example

- Effect of term frequencies

Frequency of “president”	Frequency of “lincoln”	BM25 score
15	25	20.66
15	1	12.74
15	0	5.00
1	25	18.2
0	25	15.66

# Language Model

- Language Models (LMs) assign a probability distribution over sequences of words.
- Unigram language model
  - probability distribution over the words in a language
  - generation of text consists of pulling words out of a “bucket” according to the probability distribution and replacing them
- N-gram language model
  - some applications use bigram and trigram language models where probabilities depend on previous words
- LMs are popular in a variety of applications, such as:
  - Speech recognition, machine translation, and handwriting recognition

# Language Modeling Approach

- Probability distribution over strings of text
  - How likely is a given string (observation) in a given “language”
  - For example, consider probability for the following four strings

$$p_1 > p_2 > p_3 > p_4$$

$$P1 = P(\text{“a quick brown dog”})$$

$$P2 = P(\text{“dog quick a brown”})$$


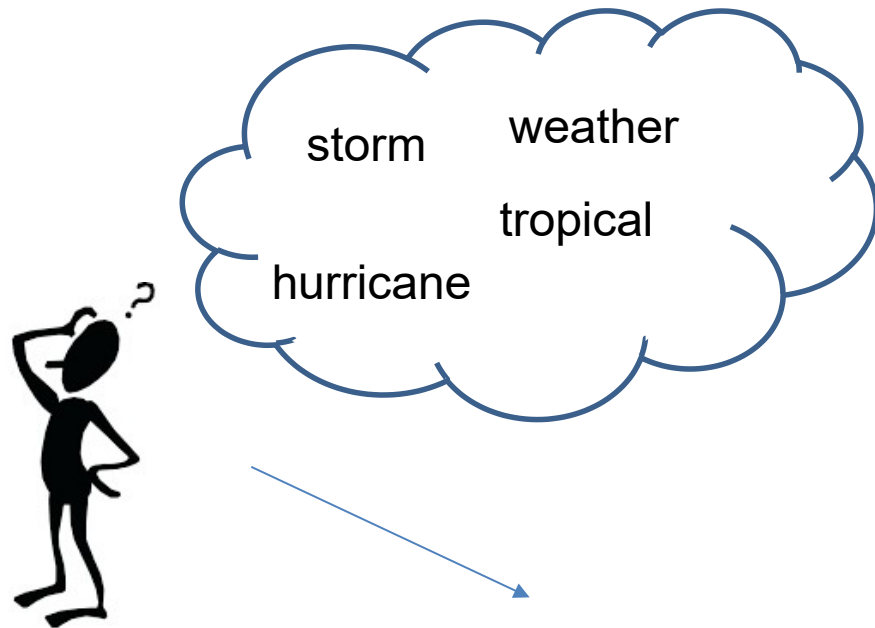
$$P3 = P(\text{“un chien quick brown”})$$

$$P4 = P(\text{“un chien brun rapide”})$$

- In IR, most likely  $P_1 = P2$

# How can we use LMs in IR?

- Use LMs to model the process of query generation:
  - User thinks of some relevant document
  - Picks some keywords to use as the query





## Relevant Docs

Forecasters are watching two tropical storms that could pose hurricane threats to the southern United States. One is a downgraded




# Retrieval with Language Model

- Each document  $D$  in the collection defines a “language”
  - All possible sentences the author of  $D$  could have written
  - $P(S|MD)$  is the probability that author would write string  $s$ 
    - Intuition: write a billion variants of  $D$ , count how many times we get “ $s$ ”
    - Language model of what the author of  $D$  was trying to say
- Retrieval: rank documents by  $P(q|MD)$ 
  - Probability that the author would write “ $q$ ” while creating  $D$



# Unigram Language Models


- Words are "sampled" independently of each other
  - Metaphor: randomly pulling out words from a jar (with replacement)
  - Joint probability decomposes into a product of marginals
  - Estimation of probabilities: simple counting



$$P(\text{red, yellow, red, blue}) = P(\text{red}) P(\text{yellow}) P(\text{red}) P(\text{blue})$$

# Estimation of Language Models


- Usually we don't know the model  $M$ 
  - But we may have a sample of text representative of that model
  - Estimate a language model from that sample
- Maximum likelihood estimator:
  - Count relative frequency of each word



$$\begin{aligned}P(\text{blue}) &= 1/3 \\P(\text{red}) &= 1/3 \\P(\text{yellow}) &= 1/3 \\P(\text{green}) &= 0 \\P(\text{grey}) &= 0\end{aligned}$$


# The zero-frequency problem

- Suppose some event not in our example
  - Language is sparse, so not every event will be present in our document.
  - Model will assign zero probability to that event (and any set of events involving the unseen event)
- It is incorrect to infer zero probabilities
  - Especially when dealing with incomplete samples



# Simple Discounting Methods


- Laplace correction:
  - Add 1 to every count, normalize
  - Problematic for large vocabularies
- Lindstone correction:
  - Add a small constant epsilon to every count
- Absolute Discounting
  - Subtract a constant epsilon from non-zero values, and re-distribute the probability mass



$$\begin{aligned}P(\text{blue}) &= (1 + \epsilon) / (3 + 5\epsilon) \\P(\text{red}) &= (1 + \epsilon) / (3 + 5\epsilon) \\P(\text{yellow}) &= (1 + \epsilon) / (3 + 5\epsilon) \\P(\text{green}) &= (0 + \epsilon) / (3 + 5\epsilon) \\P(\text{grey}) &= (0 + \epsilon) / (3 + 5\epsilon)\end{aligned}$$

# How to set Epsilon?

- How to set the discounting parameter epsilon
- Leave-one-out discounting
  - Remove some word  $w$ , compute  $P(D \mid D_{-w})$
  - Repeat for every word in the document
  - Iteratively adjust epsilon to maximize  $P(D \mid D-w)$ 
    - Intuitively: increase if word  $w$  occurs once, decrease if more than once



Step 1 - as is epsilon is zero


Step 2 – drop blue from document, then epsilon need to be updated to a value  $> 0$

Step 3 –drop red from document, since we have two reds, then epsilon will need to be updated again

...


# Interpolation Methods

- Problem with all discounting methods:
  - Discounting treats unseen words equally (add or subtract epsilon)
  - Some words are more frequent than others
- Idea – instead of a fixed epsilon for all terms, lets use probability of terms in the collection
  - Plays a similar role as IDF measure



# Jelinek-Mercer Smoothing

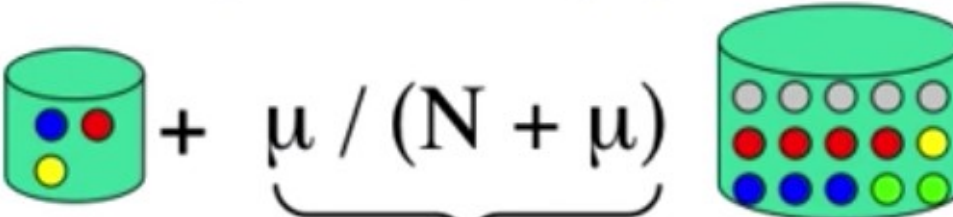
- Correctly setting lamda is very important
- Start simple
  - Set lambda to be a constant (independent of the document and query)
- Tune
  - Tune to optimize retrieval performance






# Dirichlet Smoothing

- Problem with Jelinek-Mercer
  - Issue - Using the same smoothing for documents of different length
  - Hence, longer documents will provide better estimates
- Make smoothing depend of sample size
- Formal derivation from Bayesian (Dirichlet) prior on LMs
- Best out-of-the-box choice for short queries
  - Tune  $\mu$  to optimize effectiveness

$$\underbrace{N / (N + \mu)}_{\lambda} \text{ } \text{ } + \text{ } \underbrace{\mu / (N + \mu)}_{(1-\lambda)}$$


## Role of smoothing as IDF



$$P(Q | D) = \prod_{w \in Q} \left( \lambda \frac{tf_{w,D}}{|D|} + (1 - \lambda) \frac{cf_w}{|C|} \right)$$

Probability of drawing a query Q from model D

Rank documents using this probability

# Major issues in applying LMs

- What kind of language model should we use?
  - Unigram or higher-order models?
  - Multinomial or multiple-Bernoulli?
- How can we estimate model parameters?
  - Maximum likelihood and zero frequency problem
  - Discounting methods: Laplace, Lindstone, ...
  - Interpolation methods: Jelinek-Mercer, Dirichlet prior
  - Leave-one-out method
- Ranking methods
  - Query likelihood, document likelihood, model comparison

# Why unigram models?

- Unigram = word independence
  - Best for information retrieval
- What about higher-order models?
  - N-gram models – critical in other fields like speech recognition
  - But don't really seem to work out well in IR
  - Must produce grammatical sequences of words that match documents and queries, and the probabilities are low