

# CLUSTERING


---

# Outline

- Clustering
  - Chapter 7 and 12 in Mining of Massive Datasets Textbook
  - Chapter 9 in Information Retrieval in Practice
- Today
  - General clustering (Kmeans)
- Friday
  - LDA – Latent Dirichlet Allocation



# The Problem of Clustering

- Clustering is a technique for finding similarity groups in data, called clusters.
  - Groups data instances that are
    - Similar to (near) each other in one cluster, and
    - Very different (far away) from each other into different clusters.
- Clustering is an unsupervised learning task as no class values denoting an a priori grouping of the data instances are given, which is the case in supervised learning.



# Unsupervised learning: clustering

Raw data



extract  
features

Features


$f_1, f_2, f_3, \dots, f_n$

$f_1, f_2, f_3, \dots, f_n$


$f_1, f_2, f_3, \dots, f_n$

$f_1, f_2, f_3, \dots, f_n$

$f_1, f_2, f_3, \dots, f_n$




group into  
classes/clusters



**No** “supervision”, we’re only given data and want to find natural groupings

# What do we need for clustering?

- Proximity measure, either
  - Similarity measure  $s(x_1, x_2)$  - large if  $x_1, x_2$  are similar
  - Dissimilarity (or distance) measure  $d(x_1, x_2)$  – small if  $x_1, x_2$  are similar
- Criterion function to evaluate a clustering



- Algorithm to compute clustering
  - For example, by optimizing the criterion function

# Distance (dissimilarity) measures

- **Euclidean distance**

- $d(x_i, x_j) = \sqrt{\sum_{k=1}^d (x_i^{(k)} - x_j^{(k)})^2}$

- **Manhattan (city block) distance**

- Approximation to Euclidean distance, cheaper to compute

- $d(x_i, x_j) = \sum_{k=1}^d |x_i^{(k)} - x_j^{(k)}|$

- **Minkowski distance** (there are special cases)

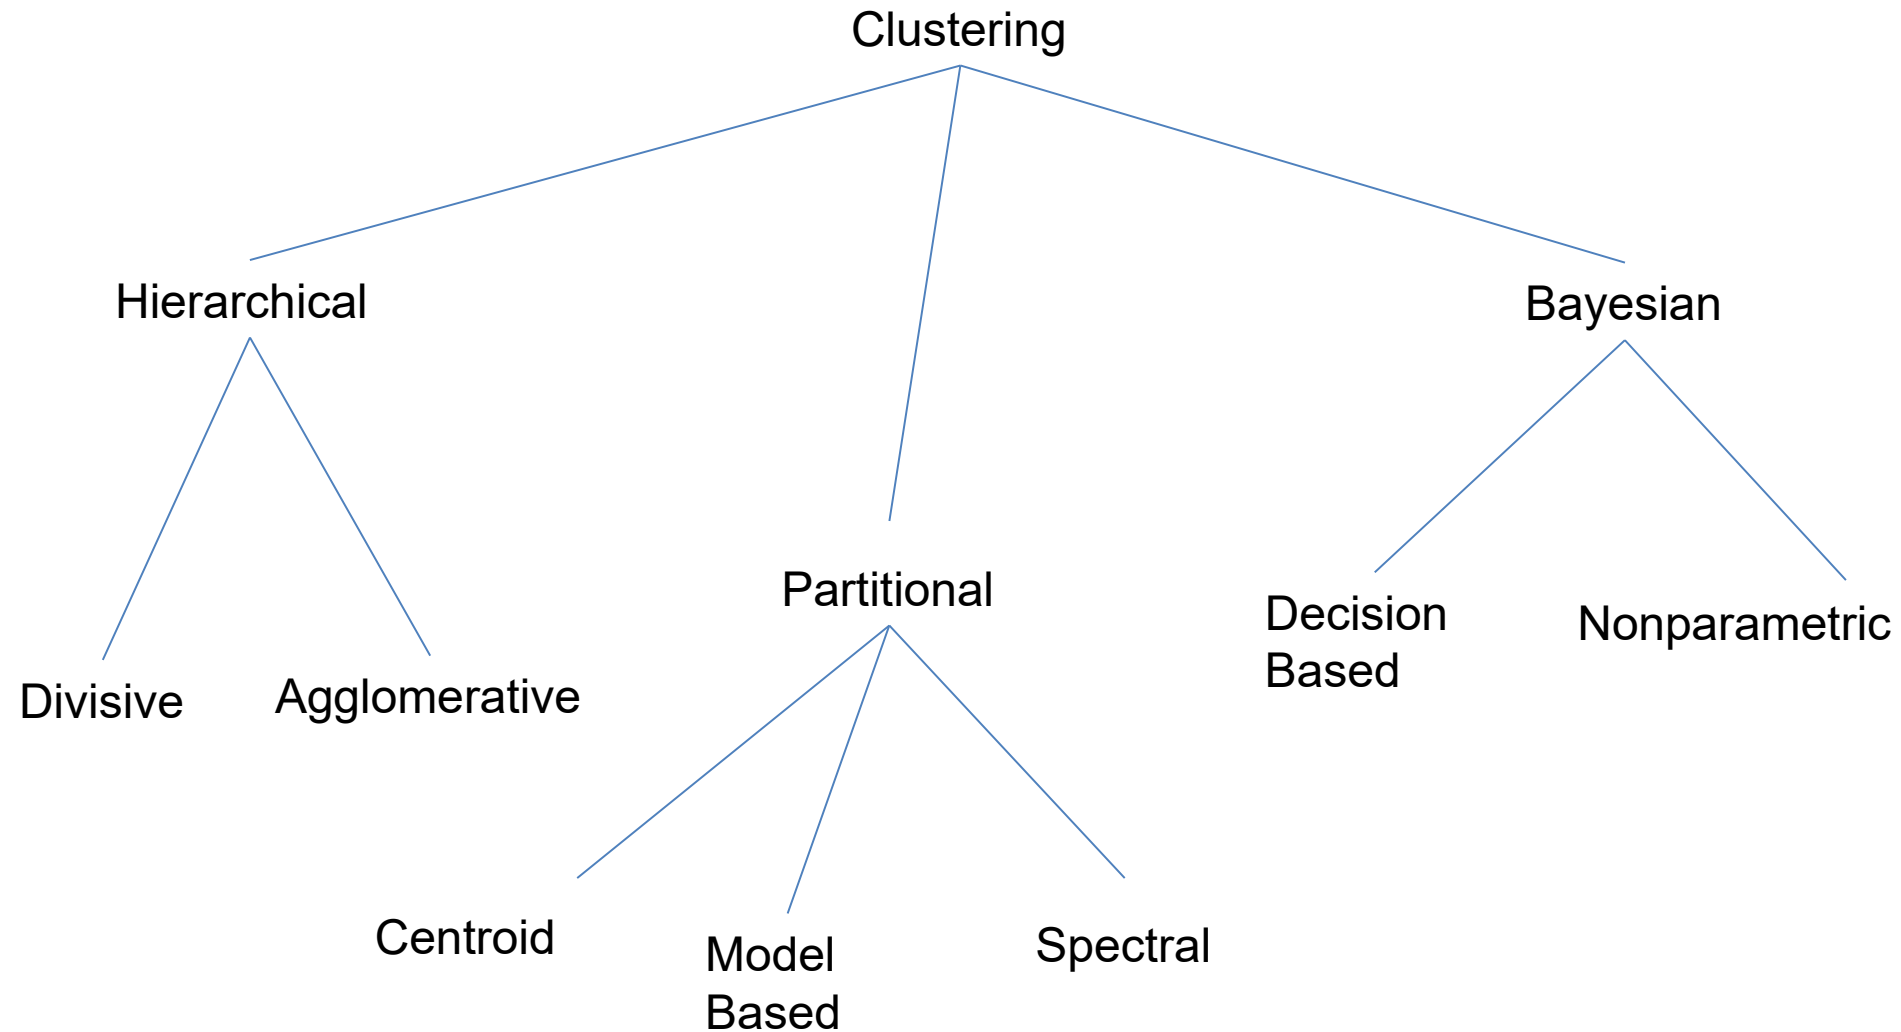
- $d_p(x_i, x_j) = (\sum_{k=1}^m |x_{ik} - x_{jk}|^p)^{\frac{1}{p}}$

- P is a positive integer

# Cluster evaluation (a hard problem)

- **Intra-cluster cohesion**(compactness):
  - Cohesion measures how near the data points in a cluster are to the cluster centroid.
  - Sum of squared error (SSE) is a commonly used measure.
- **Inter-cluster separation**(isolation):
  - Separation means that different cluster centroids should be far away from one another.
- In most applications, expert judgments are still the key

# Clustering Techniques

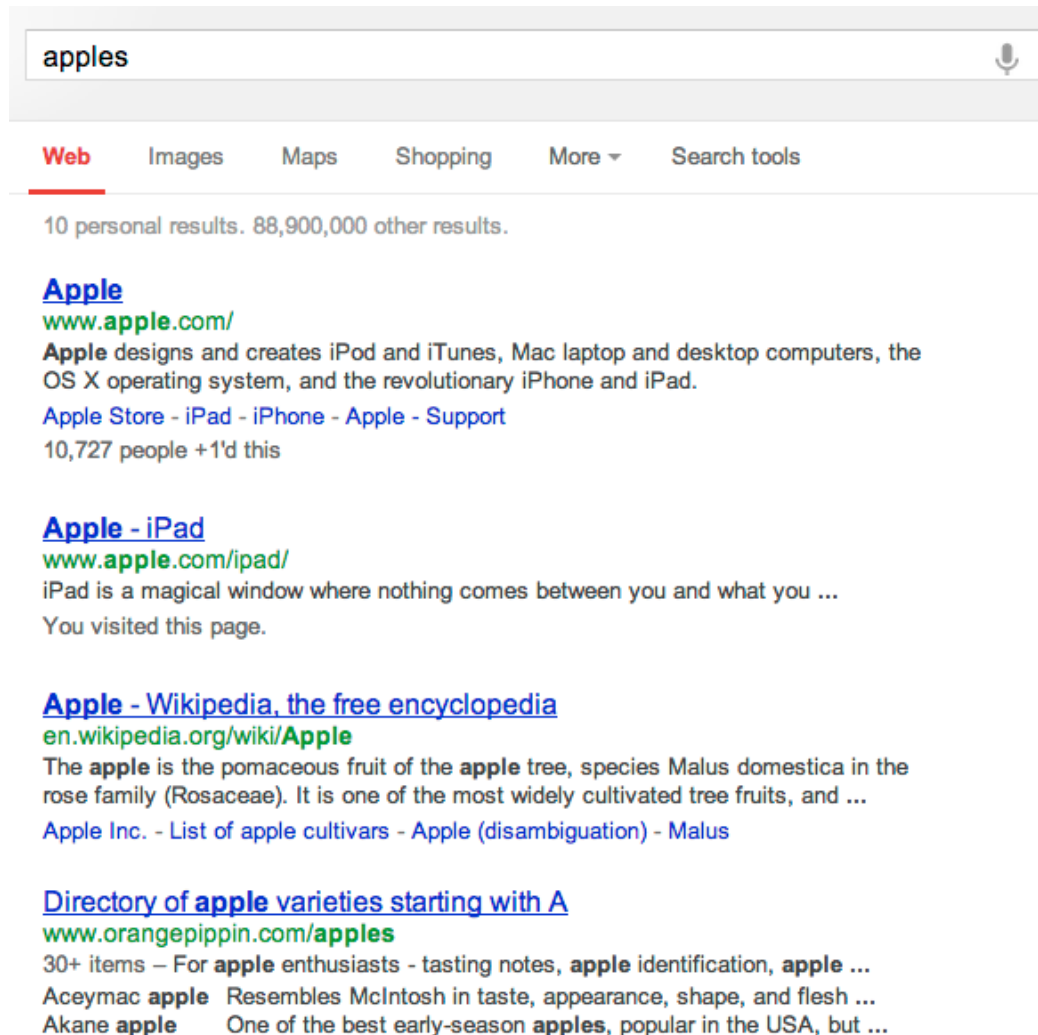




# Applications of clustering

- *Understanding*
  - Grouping objects into conceptually meaningful classes is an important step in analysis.
  - EX: Group related documents for browsing, group genes and proteins that have similar functionality, or group stocks with similar price fluctuations.
- *Summarization*
  - Reduce the size of large data sets
  - EX: PCA requires space complexity  $O(m^2)$  so may not be practice for large datasets. We can cluster dataset and apply PCA on the cluster prototypes.

# Applications (clustering for IR)



Documents or webpages  
in the same cluster are  
likely to be similar.


# Application – Clustering for image search



Given a collection of unlabeled objects(images), cluster them into groups. So that given an image, we can pull similar images from the right cluster.

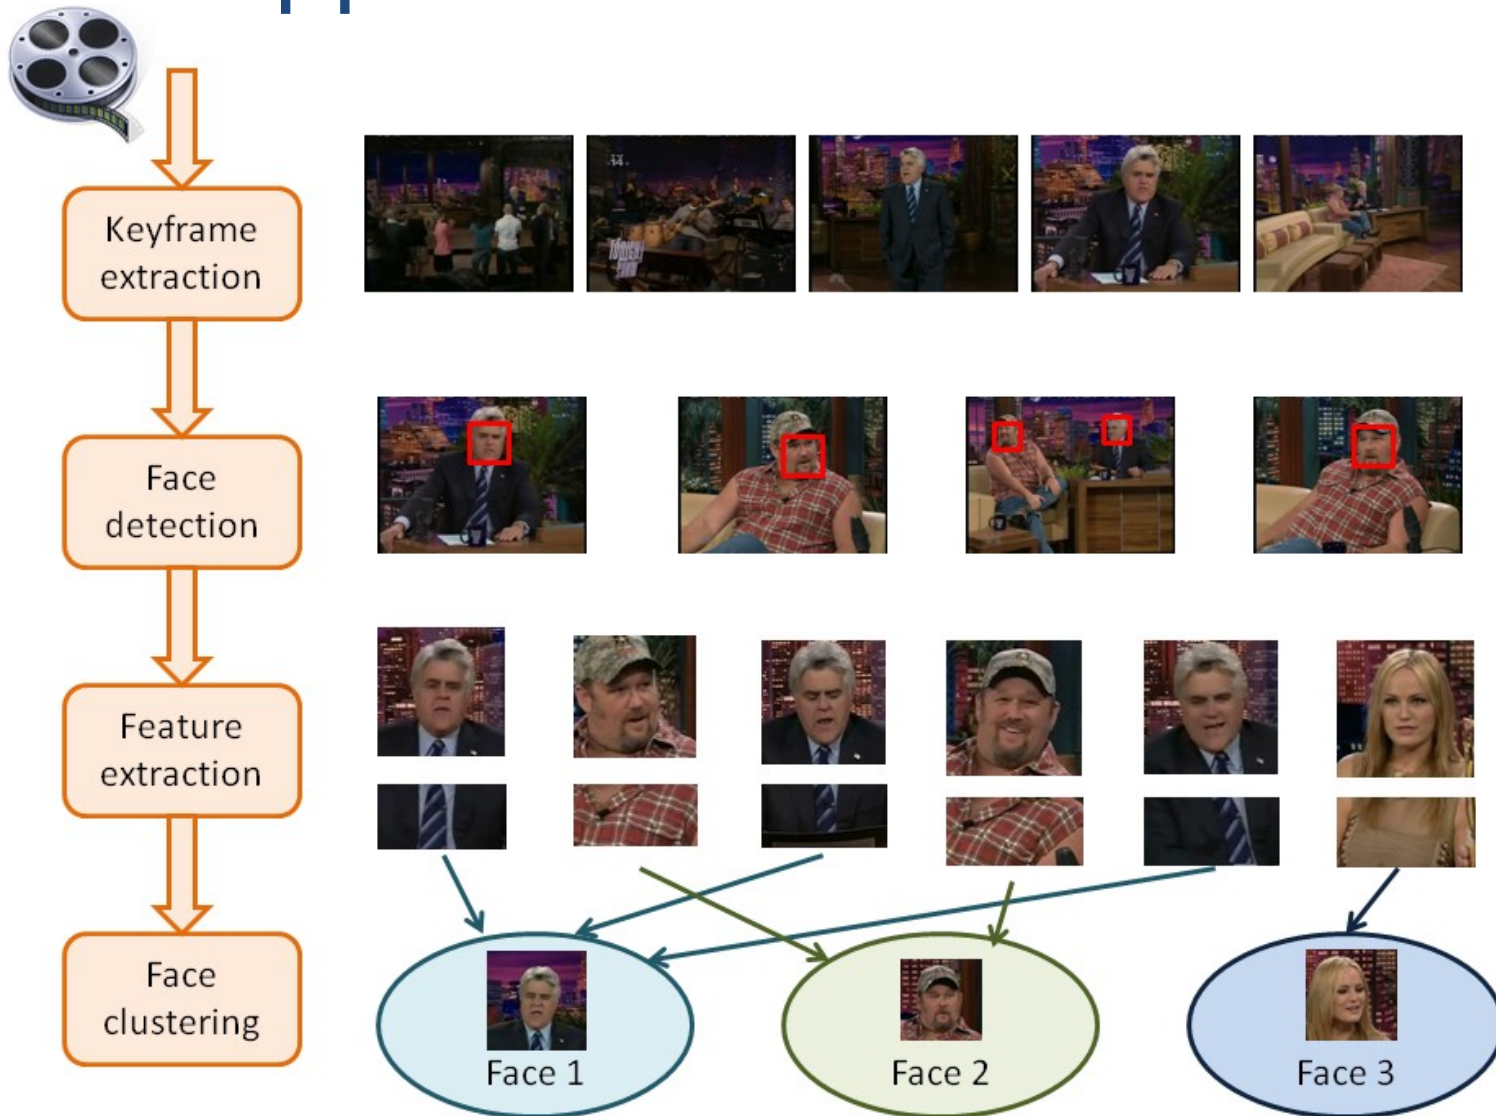
# Gene expression data

Data from Garber et al.  
PNAS (98), 2001.

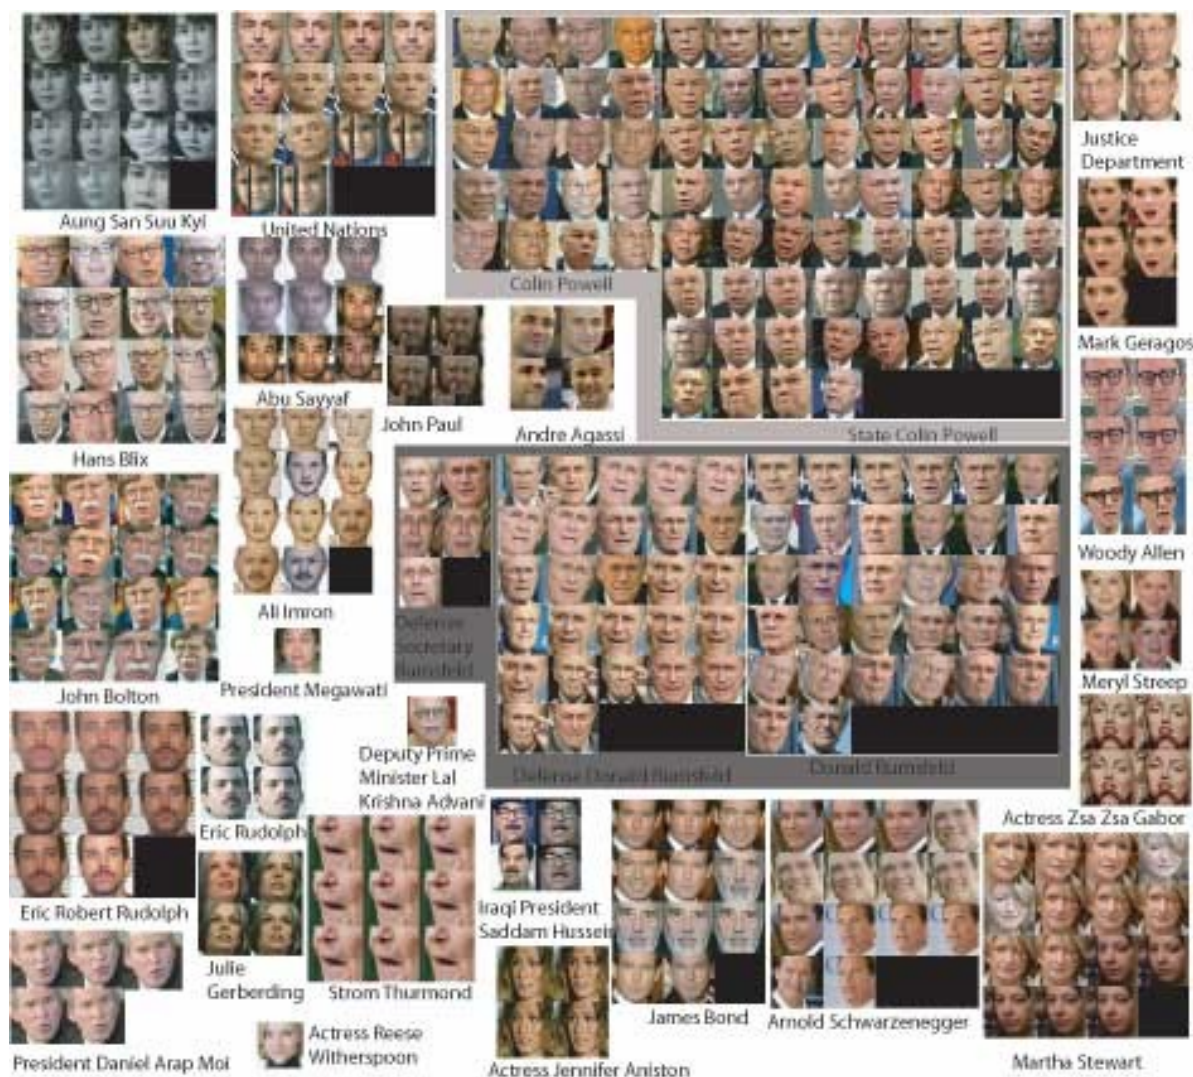




# Applications - Face Clustering




# Applications - Face clustering





# Clustering is a hard problem!



# Why is it hard?

- Clustering in two dimensions looks easy
- Clustering small amounts of data looks easy
- And in most cases, looks are not deceiving
- Many applications involve not 2, but 10 or 10,000 dimensions
- High-dimensional spaces look different: Almost all pairs of points are at about the same distance



# Overview: Methods of Clustering

- **Hierarchical:**

- **Agglomerative** (bottom up):

- Initially, each point is a cluster
    - Repeatedly combine the two “nearest” clusters into one

- **Divisive** (top down):

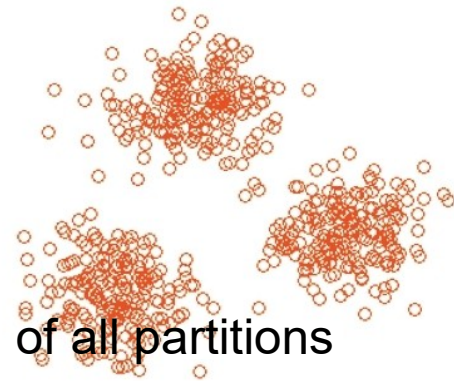

- Start with one cluster and recursively split it

- **Partitional**

- Usually start with a random (partial) partitioning
  - Refine it iteratively
    - $K$  means clustering
    - Model based clustering

- **Bayesian**

- Try to generate a posteriori distribution over the collection of all partitions of the data.



# Hard vs. soft clustering


***Hard clustering:*** Each example belongs to exactly one cluster

***Soft clustering:*** An example can belong to more than one cluster (probabilistic)

- Makes more sense for applications like creating browsable hierarchies
- You may want to put a pair of sneakers in two clusters: (i) sports apparel and (ii) shoes

# Hierarchical Clustering

- Key operation: Repeatedly combine two nearest clusters
- Three important questions:
  - 1) How do you represent a cluster of more than one point?
  - 2) How do you determine the “nearness” of clusters?
  - 3) When to stop combining clusters?



# Hierarchical Clustering

- Key operation: Repeatedly combine two nearest clusters


## (1) How to represent a cluster of many points?

- Key problem: As you merge clusters, how do you represent the “location” of each cluster, to tell which pair of clusters is closest
- Euclidean case: each cluster has a centroid = average of its (data)points

## (2) How to determine “nearness” of clusters?

- Measure cluster distances by distances of centroids


# Example: Hierarchical clustering



**Data:**

o ... data point

x ... centroid



**Dendrogram**


# And in the Non-Euclidean Case?

- What about the Non-Euclidean case?
  - The only “locations” we can talk about are the points themselves
    - i.e., there is no “average” of two points
- Approach 1:
  - (1) How to represent a cluster of many points?  
clustroid = (data)point “closest” to other points
  - (2) How do you determine the “nearness” of clusters? Treat clustroid as if it were centroid, when computing inter-cluster distances

# “Closest” Point?

- (1) How to represent a cluster of many points?  
clustroid = point “closest” to other points
- Possible meanings of “closest”:
  - Smallest maximum distance to other points
  - Smallest average distance to other points
  - Smallest sum of squares of distances to other points
    - For distance metric  $d$  clustroid  $c$  of cluster  $C$  is:

$$\min_c \sum_{x \in C} d(x, c)^2$$



**Centroid** is the avg. of all (data)points in the cluster. This means centroid is an “artificial” point.

**Clustroid** is an **existing** (data)point that is “closest” to all other points in the cluster.

# Implementation

- Naïve implementation of hierarchical clustering:
  - At each step, compute pairwise distances between all pairs of clusters, then merge
  - $O(N^3)$
- Careful implementation using priority queue can reduce time to  $O(N^2 \log N)$ 
  - Still too expensive for really big datasets that do not fit in memory



# K-MEANS CLUSTERING

---

Demo of K-means clustering :

<http://www.onmyphd.com/?p=kmeans.clustering&ckattempt=2>


# K-means

- Given  $k$ , the *k-means* algorithm works as follows:
  1. Choose  $k$ (random) data points (**seeds**) to be the initial **centroids**, cluster centers
  2. Assign each data point to the closest **centroid**
  3. Re-compute the **centroids** using the current cluster memberships
  4. If a convergence criterion is not met, repeat steps 2 and 3

# K-means convergence (stopping) criterion

- No (or minimum) re-assignments of data points to different clusters, *or*
- No (or minimum) change of centroids, *or*
- Minimum decrease in the **sum of squared error**(SSE),
  - $SSE = \sum_{j=1}^k \sum_{x \in C_j} d(x, m_j)^2$
  - $C_j$  is the  $j^{\text{th}}$  cluster,
  - $\mathbf{m}_j$  is the centroid of cluster  $C_j$  (the mean vector of all the data points in  $C_j$ ),
  - $d(\mathbf{x}, \mathbf{m}_j)$  is the (Euclidian) distance between data point  $\mathbf{x}$  and centroid  $\mathbf{m}_j$


# K-means: an example



## Randomly Initialize the cluster centers

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster




## Iteration # 1


- **Assign each item to closest cluster center**

iterate over each point:

- get distance to each cluster center
  - assign to closest center (hard cluster)
- Recalculate centers as the mean of the points in a cluster




After assigning points to nearest center



## Iteration # 1

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster





# Distance measures

Euclidean:


$$d(x, y) = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

good for spatial data

# Clustering documents

One feature for each word. The value is the number of times that word occurs.


Documents are points or vectors in this space



# When Euclidean distance doesn't work


Which document is closest to  $q$  using Euclidean distance?

Which do you think should be closer?



# Issues with Euclidian distance




- the Euclidean distance between  $q$  and  $d_2$  is large
- but, the distribution of terms in the query  $q$  and the distribution of terms in the document  $d_2$  are very similar
- This is not what we want!



# cosine similarity

$$\text{sim}(x, y) = \frac{x \bullet y}{|x||y|} = \frac{x}{|x|} \bullet \frac{y}{|y|} = \frac{\sum_{i=1}^n x_i y_i}{\sqrt{\sum_{i=1}^n x_i^2} \sqrt{\sum_{i=1}^n y_i^2}}$$

correlated with the  
angle between two  
vectors



# cosine distance

cosine similarity is a similarity between 0 and 1, with things that are similar 1 and not 0

We want a distance measure, cosine distance:


$$d(x, y) = 1 - \text{sim}(x, y)$$

- good for text data and many other “real world” data sets
- is computationally friendly since we only need to consider features that have non-zero values in **both** examples

# K-means

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster



How do we calculate these?

# K-means

Iterate:


- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

Mean of the points in the cluster:

$$\mu(C) = \frac{1}{|C|} \sum_{x \in C} x$$

where:

$$x + y = \sum_{i=1}^n x_i + y_i \quad \frac{x}{|C|} = \sum_{i=1}^n \frac{x_i}{|C|}$$





# Comments on the K-Means Method

- Strength

- Relatively efficient:  $O(tkn)$ ,
  - where  $n$  is # objects,
  - $k$  is # clusters, and
  - $t$  is # iterations.
  - Normally,  $k, t \ll n$ .
- Often terminates at a local optimum.

- Weakness

- Applicable only when mean is defined, then what about categorical data?
- Need to specify  $k$ , the number of clusters, in advance
- Unable to handle noisy data and outliers
- Not suitable to discover clusters with non-convex shapes

# K-means loss function

K-means tries to minimize what is called the “k-means” loss function:

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

that is, the sum of the squared distances from each point to the associated cluster center

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

Does each step of k-means move towards reducing this loss function (or at least not increasing)?

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

Intuition:

1. Any other assignment would end up in a larger loss
2. The mean of a set of values minimizes the squared error

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

Does this mean that k-means will always find the minimum loss/clustering?

# Minimizing k-means loss

Iterate:

1. Assign/cluster each example to closest center
  2. Recalculate centers as the mean of the points in a cluster
- 

$$loss = \sum_{i=1}^n d(x_i, \mu_k)^2 \quad \text{where } \mu_k \text{ is cluster center for } x_i$$

NO! It will find *a minimum*.

Unfortunately, the k-means loss function is generally not convex and for most problems has many, many minimums

We're only guaranteed to find one of them

# K-means variations/parameters

Start with some initial cluster centers

Iterate:

- Assign/cluster each example to closest center
- Recalculate centers as the mean of the points in a cluster

What are some other variations/parameters we haven't specified?

# K-means variations/parameters

Initial (seed) cluster centers


## Convergence

- A fixed number of iterations
- partitions unchanged
- Cluster centers don't change

K!



# K-means: Initialize centers randomly



What would happen  
here?

Seed selection  
ideas?

# Seed choice

Results can vary drastically based on random seed selection

Some seeds can result in poor convergence rate, or convergence to sub-optimal clusterings

## Common heuristics

- Random centers in the space
- Randomly pick examples
- Points least similar to any existing center (furthest centers heuristic)
- **Try out multiple starting points**
- Initialize with the results of another clustering method

# Furthest centers heuristic

$\mu_1$  = pick random point

for  $i = 2$  to  $K$ :

$\mu_i$  = point that is furthest from **any** previous centers


---

$$\mu_i = \underset{x}{\operatorname{argmax}} \underbrace{\min_{\mu_j : 1 < j < i} d(x, \mu_j)}_{\text{smallest distance from } x \text{ to any previous center}}$$

point with the largest  
distance to any previous  
center


smallest distance from  $x$  to  
any previous center

# K-means: Initialize furthest from centers




Pick a random point for the first center

## K-means: Initialize furthest from centers



What point will be chosen next?


# K-means: Initialize furthest from centers



Furthest point from  
center

What point will be chosen  
next?


# K-means: Initialize furthest from centers



Furthest point from  
center

What point will be chosen  
next?

# K-means: Initialize furthest from centers




Furthest point from  
center

Any issues/concerns with this  
approach?




# Furthest points concerns




If  $k = 4$ , which points will get chosen?

# Furthest points concerns



If we do a number of trials, will we get different centers?

# Furthest points concerns



Doesn't deal well with outliers

# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to **K**:

for  $i = 1$  to **N**:

$s_i = \min d(x_i, \mu_{1\dots k-1})$  // smallest distance to any center

$\mu_k$  = randomly pick point *proportionate* to *s*

How does this  
help?

# K-means++

$\mu_1$  = pick random point

for  $k = 2$  to **K**:


for  $i = 1$  to **N**:

$s_i = \min d(x_i, \mu_{1\dots k-1})$  // smallest distance to any center

- $\mu_k$  = randomly pick point *proportionate* to **s**
  - Makes it possible to select other points
  - if #points >> #outliers, we will pick good points
- Makes it non-deterministic, which will help with random runs
- Nice theoretical guarantees!


# Getting the $k$ right

- How to select  $k$ ?
- Try different  $k$ , looking at the change in the average distance to centroid as  $k$  increases
- Average falls rapidly until right  $k$ , then changes little




# Example: Picking $k$

**Too few;**  
many long  
distances  
to centroid.



# Example: Picking $k$


**Just right;**  
distances  
rather short.





# Example: Picking $k$

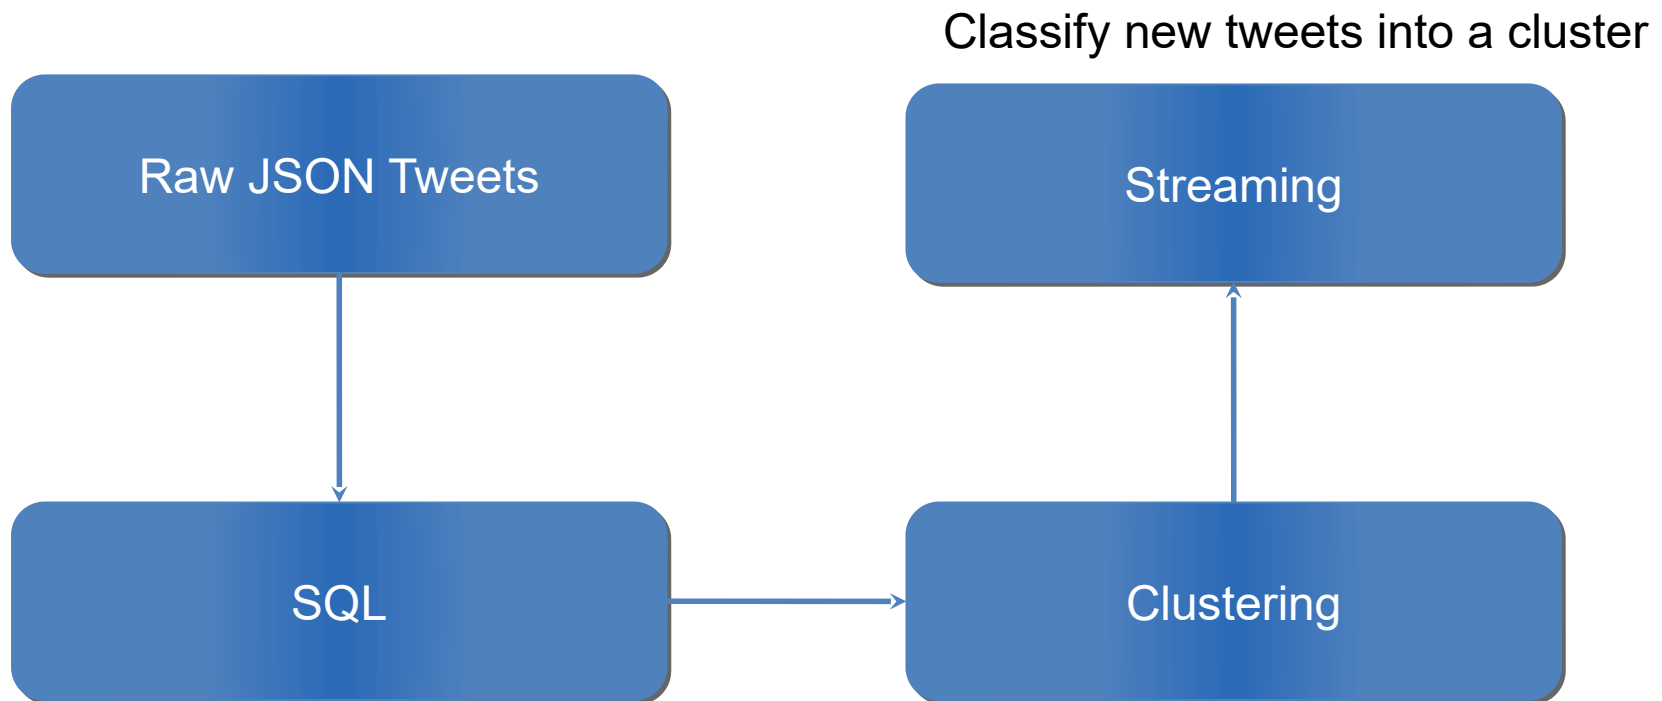
**Too many;**  
little improvement  
in average  
distance.



# Tweet Clustering by Language

- Tutorial
  - <http://zdatainc.com/2014/08/real-time-streaming-apache-spark-streaming/>
- Code is on github:
  - <https://github.com/databricks/reference-apps>
  - Uses Scala instead of Python (note, Twitter Stream is not yet natively supported for Python).

# Cluster Tweets



Load raw tweets into table,  
then use SQL to pull out  
relevant information to feed  
into clustering algorithm.

# Steps

- The problem is broken up into 3 parts :
- **Collect a Dataset of Tweets** - Spark Streaming is used to collect a dataset of tweets and write them out to files.
- **Examine the Tweets and Train a Model** - Spark SQL is used to examine the dataset of Tweets. Then Spark MLLib is used to apply the K-Means algorithm to train a model on the data.
- **Apply the Model in Real-time** - Spark Streaming and Spark MLLib are used to filter a live stream of Tweets for those that match the specified cluster.

# Collecting Dataset

- [Step 1](#) – collect dataset

```
val tweetStream = TwitterUtils.createStream(ssc, Utils.getAuth)  
    .map(gson.toJson(_))
```

```
tweetStream.foreachRDD((rdd, time) => {  
    val count = rdd.count()  
    if (count > 0) {  
        val outputRDD = rdd.repartition(partitionsEachInterval)  
        outputRDD.saveAsTextFile(  
            outputDirectory + "/tweets_" + time.milliseconds.toString)  
        numTweetsCollected += count  
        if (numTweetsCollected > numTweetsToCollect) {  
            System.exit(0)  
        }  
    }  
})
```

# Cluster Tweets

## Step 2 – Cluster tweets

```
val vectors = texts.map(Utils.featurize).cache()
  vectors.count() // Calls an action on the RDD to populate the vectors cache.
  val model = KMeans.train(vectors, numClusters, numIterations)
  sc.makeRDD(model.clusterCenters,
numClusters).saveAsObjectFile(outputModelDir)
```

```
val some_tweets = texts.take(100)
println("----Example tweets from the clusters")
for (i <- 0 until numClusters) {
  println(s"\nCLUSTER $i:")
  some_tweets.foreach { t =>
    if (model.predict(Utils.featurize(t)) == i) {
      println(t)
    }
  }
}
```

# Classify New Tweets

- [Step 3](#) – Classify new tweets

```
println("Initializing Streaming Spark Context...")  
val conf = new SparkConf().setAppName(this.getClass.getSimpleName)  
val ssc = new StreamingContext(conf, Seconds(5))
```

```
println("Initializing Twitter stream...")  
val tweets = TwitterUtils.createStream(ssc, Utils.getAuth)  
val statuses = tweets.map(_.getText)
```

```
println("Initializing the KMeans model...")  
val model = new KMeansModel(ssc.sparkContext.objectFile[Vector](  
    modelFile.toString).collect())
```

```
val filteredTweets = statuses  
    .filter(t => model.predict(Utils.featurize(t)) == clusterNumber)  
filteredTweets.print()
```

# Latent Dirichlet Allocation (LDA)

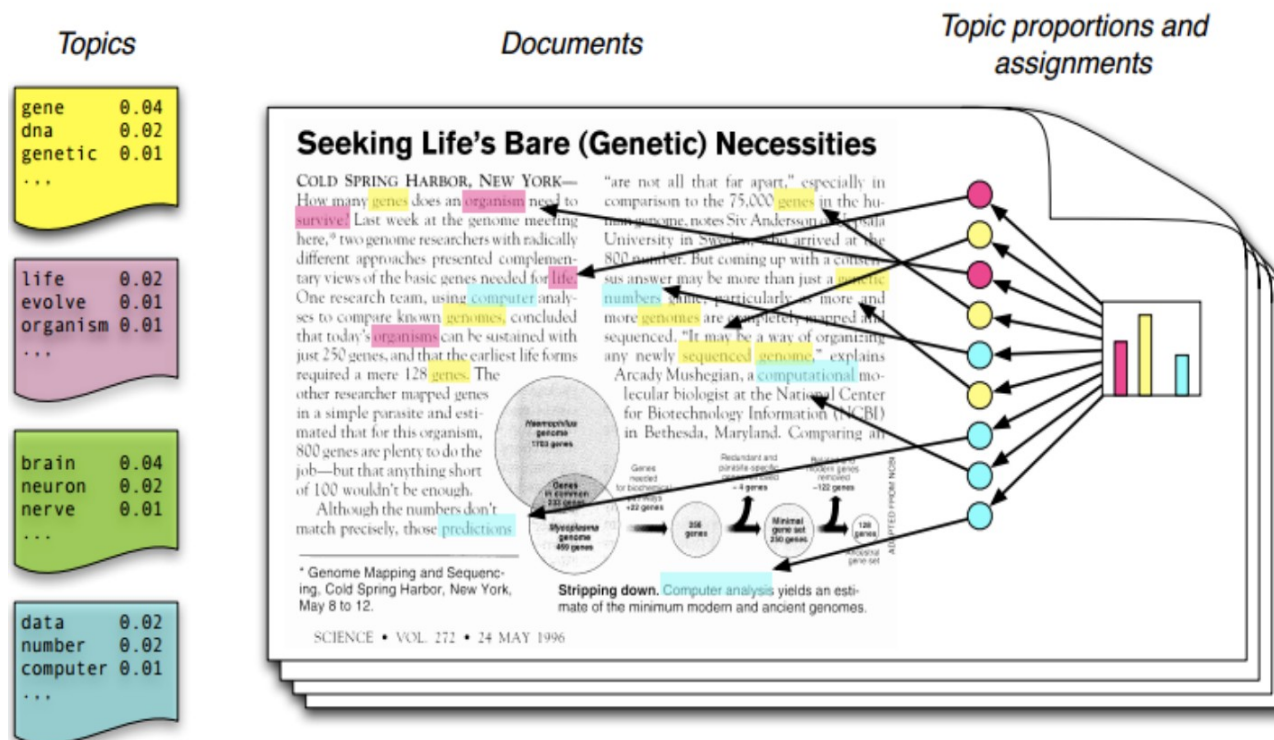
- As more information becomes available, it becomes more difficult to find and discover what we need.
- We need tools to help us organize, search and understand these vast amount of information.
- Topic modeling provides methods for automatically organizing, understanding, searching, and summarizing large electronic archives:
  - 1. Discover the hidden themes in the collection
    2. Annotate the documents according to these themes
    3. Use annotations to organize, summarize, search, and form predictions



# Some Assumptions

- We have a set of documents  $D_1, D_2, \dots, D_n$ .
- Each document is just a collection of words or a “bag of words”. Thus, the order of the words and the grammatical role of the words (subject, object, verbs, ...) are not considered in the model.
- Stop words like (am, is, the, this...) can be eliminated from the documents as a preprocessing step since they don't carry any information about the “topics”.
- In fact, we can eliminate words that occur in at least %80 ~ %90 of the documents!
- Each document is composed of  $N$  “important” or “effective” words, and we want  $K$  topics.

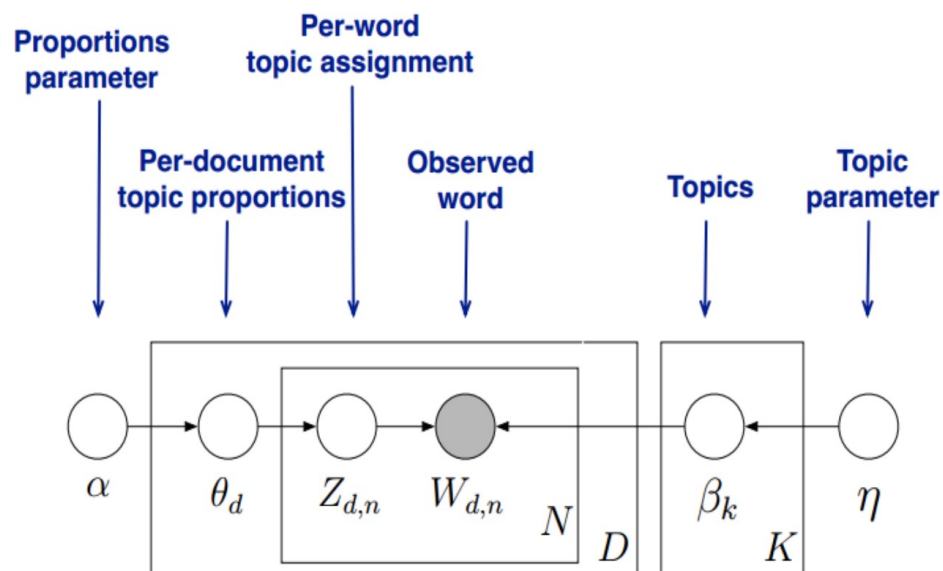
# Model Definition



- Each topic is a distribution over words
- Each document is a mixture of corpus-wide topics
- Each word is drawn from one of these topics
- We only observe the words within the documents and the other structure are hidden variables.

# Model Definition

- Our goal is to infer or estimate the hidden variables, i.e. computing their distribution conditioned on the documents.
  - $P(\text{topics, proportions, assignments} \mid \text{documents})$
- Plate notation is a way of visually representing the dependencies among the model parameters.



- Shaded nodes are observed, and unshaded nodes are hidden.
- $\alpha$  – is the parameter of Dirichlet prior on the per-document topic distribution
- $\beta$  – is the parameter of the Dirichlet prior on the per-topic word distribution
- $\mathbf{W}_{d,n}$  – a specific word
- $\mathbf{Z}_{d,n}$  – is the topic of the  $n$ th word in the document
- $\theta_d$  is the topic distribution for document  $d$

# Generative Process

- LDA assumes that new documents are created in the following way:
  - Determine number of words in the document
  - Choose a topic mixture for the document over a fixed set of topics (i.e. 20% topic A, 30% topic B, 50% topic C)
  - Generate the words in the documents by:
    - First pick a topic based on the documents' multinomial distribution above.
    - Next pick a word based on the topic's multinomial distribution

# Generative Process Example

- Say we have a group of articles that can be characterized by three topics: Animals, Cooking and Politics.
- Each of those topics can be described by the following words:
  - Animals: dog, chicken, cat, zoo, giraff
  - Cooking: Oven, food, restaurant, plates, taste, delicious
  - Politics: Republican, Democrat, Congress, Divisive, Ineffective
- Say we want to generate a new document that is 80% about animals and 20% about cooking.
  - We choose the length of the article (say, 1000 words)
  - We choose a topic based on our specified mixture (so, out of our 1000 words, 800 will be from topic 'animals')
  - We choose a word based on the word distribution of each topic

# Working Backwards

- Suppose you have a corpus of documents
- You want LDA to learn the topic representation of  $K$  topics in each document and the word distribution of each topic
- LDA backtracks from the document level to identify topics that are likely to have generated the corpus.

# Working Backwards (Cont.)

- Randomly assign each word in each document to one of the  $K$  topics.
- For each document  $d$ :
  - Assume that all topic assignments except for the current one are correct.
  - Calculate two proportions:
    - 1. Proportion of words in document  $d$  that are currently assigned to topic  $t$  =  $p(\text{topic } t \mid \text{document } d)$
    - 2. Proportion of assignments to topic  $t$  over all document that come from this word  $w$  =  $p(\text{word } w \mid \text{topic } t)$
  - Multiply those two proportions and assign  $w$  a new topic based on that probability.  $P(\text{topic } t \mid \text{document } d) * p(\text{word } w \mid \text{topic } t)$